



**INSTITUTE FOR ADVANCED COMPUTING AND
SOFTWARE DEVELOPMENT (IACSD), AKURDI,
PUNE**

Documentation on

“Used Cars Price Prediction”

PG-DBDA March 2023

Submitted By:

Group no. 17

Roll No.

233546

233539

Name:

Shreya Choudhary

Rinki Darade

Mr. Abhijit Nagargoje

Project Guide

Mr. Rohit Puranik

Centre Coordinator

Abstract

Determining whether the listed price of a used car is a challenging task, due to the many factors that drive a used vehicle's price on the market. The focus of this project is developing machine learning models that can accurately predict the price of a used car based on its features, in order to make informed purchases. We implement and evaluate various learning methods on a dataset consisting of the sale prices of different makes and models across cities in the United States. Our results show that catboost model and Neural Network yield the best results, but are compute heavy. Conventional linear regression also yielded satisfactory results, with the advantage of a significantly lower training time in comparison to the aforementioned methods.

Acknowledgement

I take this occasion to thank God, almighty for blessing us with his grace and taking our endeavor to a successful culmination. I extend my sincere and heartfelt thanks to our esteemed guide, **Mr. Abhijit Nagargoje** for providing me with the right guidance and advice at the crucial juncture and for showing me the right way. I extend my sincere thanks to our respected Centre Co-Ordinator **Mr. Rohit Puranik**, for allowing us to use the facilities available. I would like to thank the other faculty members also, at this occasion. Last but not the least, I would like to thank my friends and family for the support and encouragement they have given me during the course of our work.

Shreya Choudhary (233546)

Rinki Darade (233539)

Table Of Contents

Abstract.....	1
Acknowledgement	2
1.Introduction.....	5
1.1 Problem Statement	5
1.2 Aim and Objective	5
1.3 Dataset.....	6
2. Project workflow	7
2.1 Big data ETL	7
2.2 Data cleaning and preprocessing	10
2.3 Exploratory Data Analysis	12
2.4 Feature Selection.....	15
2.5 Splitting data in X and Y.....	18
2.6 Train test split and PCA.....	18
3. Building Machine Learning Models	19
3.1 Lasso	19
3.2 Ridge.....	20
3.3 Random Forest.....	21
3.4 Catboost	21
3.5 XGBoost	23
3.6 Neural Network.....	24
3.7 Best Model?	26
3.8 Random Forest in Pyspark	26
4. Transformations on Pyspark	28
5. Analysis on Tableau.....	32
6. Conclusion	37
7. References.....	38

List of Figures

5. Analysis on Tableau	32
5.1 State map by Price filtered by Year and Make	32
5.2 Top 10 Models by Price.....	32
5.3 Bottom 10 Models by Price	33
5.4 Top 10 Make by Price.....	33
5.5 Bottom 10 Make by Price	34
5.6 Dashboard of Top and Bottom features	34
5.7 Top 10 Makes by Price pie chart	35
5.8 Price by Mileage Scatter plot	35
5.9 Price by Year bar chart	36

1. Introduction

1.1 Problem Statement

Used cars Price prediction using various Machine Learning and Deep learning Algorithms and comparing the evaluation metrics for all.

1.2 Abstract

Determining whether the listed price of a used car is a challenging task, due to the many factors that drive a used vehicle's price on the market. The focus of this project is developing machine learning models that can accurately predict the price of a used car based on its features, in order to make informed purchases. We implement and evaluate various learning methods on a dataset consisting of the sale prices of different makes and models across cities in the United States. Our results show that Random Forest model and K-Means clustering with linear regression yield the best results, but are compute heavy. Conventional linear regression also yielded satisfactory results, with the advantage of a significantly lower training time in comparison to the aforementioned methods.

1.3 Aim and objective

The objective is to model the price of used cars with the available independent variables. This model will then be used by the clients to understand how exactly the prices vary with the variables. They can accordingly manipulate the strategy of the firm and concentrate on areas that will yield high returns. Further, the model will be a good way for the clients to understand the pricing dynamics of a new market. Moreover, this might give the management team an insight of the real time world.

1.4 Dataset

For this project, we are using the dataset on used car sales from all over the United States, available on Kaggle. This dataset consists of over 8 Lakh rows and 8 columns in total. Out of all the columns, our target variable is Price of the car and The independent variables are: Year, Mileage, City, State, Vin, Make, Model.

```

✓ [4] 1 #loading the data
3a 2 ## true unprocessed data
3
4 df = pd.read_csv("/content/drive/MyDrive/pro - Shreya & Rinky/true_car_listings.csv")
5
6 df.shape

```

(852122, 8)

```

✓ [5] 1 df.sample(10)
0a

```

	Price	Year	Mileage	City	State	Vin	Make	Model	
669223	20000	2015	26525	Elk River	MN	KNMAT2MV7FP507751	Nissan	RogueSV	
788581	24008	2015	18849	Baltimore	MD	5TDKK3DCXFS639422	Toyota	SiennaLE,	
797739	16988	2008	95436	Dyersburg	TN	5TEJU62N28Z472202	Toyota	Tacoma2WD	
167163	36900	2014	35941	Newport News	VA	3GCUKSEC5EG323408	Chevrolet	Silverado	
189301	15383	2015	22750	Danvers	MA	1C3CCCBB4FN585301	Chrysler	200S	
812356	32797	2016	32609	Hendersonville	TN	JTEBU5JR6G5306197	Toyota	4Runner4WD	
50241	32998	2014	37677	Tempe	AZ	WBA1J7C59EVW84524	BMW	2	
322371	32888	2015	27516	Milwaukie	OR	1FTEW1E80FFB03272	Ford	F-1504WD	
748269	13995	2014	20995	West Allis	WI	2T1BURHE8EC175403	Toyota	CorollaS	
100008	16995	2017	15546	Matteson	IL	1G1BF5SM2H7190898	Chevrolet	CruzeSedan	

2. Overall Workflow

2.1 Big data ETL

We tried one approach where we implemented ETL on Hadoop services and made a Random forest model on Pyspark to check the accuracy it would provide.

Uploading our dataset into hdfs

```
talentum@talentum-virtual-machine:~$ hdfs dfs -mkdir -p /user/talentum
talentum@talentum-virtual-machine:~$ hdfs dfs -mkdir -p /user/talentum/cars_all
talentum@talentum-virtual-machine:~$ hdfs dfs -mkdir -p
-mkdir: Not enough arguments: expected 1 but got 0
Usage: hadoop fs [generic options] -mkdir [-p] <path> ...
talentum@talentum-virtual-machine:~$ hdfs dfs -ls
Found 1 items
drwxr-xr-x  - talentum supergroup          0 2023-08-18 21:27 cars_all
talentum@talentum-virtual-machine:~$ cd shared/
talentum@talentum-virtual-machine:~/shared$ ls
cars  cars_dataset
talentum@talentum-virtual-machine:~/shared$ cd cars_dataset/
talentum@talentum-virtual-machine:~/shared/cars_dataset$ ls
cars_clean_50k.csv  cars_final.ipynb  cars_raw_50k.csv  cars_spark.ipynb  true_car_listings.csv
talentum@talentum-virtual-machine:~/shared/cars_dataset$ hdfs dfs -put true_car_listings.csv /user/talentum/cars_all
talentum@talentum-virtual-machine:~/shared/cars_dataset$ hdfs dfs -ls /user/talentum/cars_all
Found 1 items
-rw-r--r--  1 talentum supergroup  56287873 2023-08-18 21:29 /user/talentum/cars_all/true_car_listings.csv
```

Creating a database on hive and using it to make a new table named 'cars'. We have to provide the same parameters and data types as our original data to accommodate values correctly.

```
talentum@talentum-virtual-machine:~$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/talentum/hive/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/talentum/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/home/talentum/hive/lib/hive-common-2.3.6.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> create database if not exists project;
OK
Time taken: 2.209 seconds

hive> show databases;
OK
default
project
Time taken: 0.225 seconds, Fetched: 2 row(s)
hive> use project;
OK
Time taken: 0.132 seconds
hive> drop table if exists cars purge;
OK
Time taken: 0.134 seconds
```

Created table and checked first 10 values

```
hive> create external table cars (Price INT, Year INT, Mileage INT, City String, State String, Vin String, Make String, Model String) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' LOCATION '/user/talentum/cars_all/' TBLPROPERTIES ("skip.header.line.count"="1");
OK
Time taken: 0.2 seconds
hive> select * from cars limit 10;
OK
8995 2014 35725 El Paso TX 19VDE2E53EE000083 Acura ILX6-Speed
18088 2013 19606 Long Island City NY 19VDE1F52DE012636 Acura ILX5-Speed
8995 2013 48851 El Paso TX 19VDE2E52DE000025 Acura ILX6-Speed
10999 2014 39922 Windsor CO 19VDE1F71EE003817 Acura ILX5-Speed
14799 2016 22142 London UT 19UDE2F32GA001284 Acura ILXAutomatic
7989 2012 105246 Miami FL 3H4CU2F83CC019895 Acura TSXAAutomatic
14490 2014 34032 Greatneck NY 3H4CU2F84EC002686 Acura TSXSspecial
13995 2013 32384 West Jordan UT 3H4CU2F64DC006203 Acura TSX5-Speed
10495 2013 57596 Waterbury CT 19VDE2E50DE000234 Acura ILX6-Speed
9995 2013 63887 El Paso TX 19VDE1F50DE010450 Acura ILX5-Speed
Time taken: 0.38 seconds, Fetched: 10 row(s)
hive>
```


Pyspark and Hive connectivity

```
In [1]: # Initialization
import os
import sys

os.environ["SPARK_HOME"] = "/home/talentum/spark"
os.environ["PYLIB"] = os.environ["SPARK_HOME"] + "/python/lib"
# In below two lines, use /usr/bin/python2.7 if you want to use Python 2
os.environ["PYSPARK_PYTHON"] = "/usr/bin/python3.6"
os.environ["PYSPARK_DRIVER_PYTHON"] = "/usr/bin/python3"
sys.path.insert(0, os.environ["PYLIB"] + "/py4j-0.10.7-src.zip")
sys.path.insert(0, os.environ["PYLIB"] + "/pyspark.zip")

# NOTE: Whichever package you want mention here.
# os.environ['PYSPARK_SUBMIT_ARGS'] = '--packages com.databricks:spark-xml_2.11:0.6.0 pyspark-shell'
# os.environ['PYSPARK_SUBMIT_ARGS'] = '--packages org.apache.spark:spark-avro_2.11:2.4.0 pyspark-shell'
os.environ['PYSPARK_SUBMIT_ARGS'] = '--packages com.databricks:spark-xml_2.11:0.6.0,org.apache.spark:spark-avro_2.11:2.4.0 pyspark-shell'
# os.environ['PYSPARK_SUBMIT_ARGS'] = '--packages com.databricks:spark-xml_2.11:0.6.0,org.apache.spark:spark-avro_2.11:2.4.0 pyspark-shell'
```

```
In [2]: #Entrypoint 2.x
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName("Spark SQL basic example").enableHiveSupport().getOrCreate()

# On yarn:
spark = SparkSession.builder.appName("Spark SQL basic example").enableHiveSupport().master("yarn").getOrCreate()
#specify.master("yarn")

sc = spark.sparkContext
```

```
In [3]: # Here is the code for making connectivity to hive metastore from pyspark

spark = spark.builder.master('yarn').config('spark.sql.warehouse.dir', '/user/hive/warehouse').config('hive.metastore', 'org.apache.hive.hcatalog.meterstore').getOrCreate()
spark.sql("show databases").show()
spark.sql("use project").show()
spark.sql("show tables").show()
print(" - ")
spark.table('cars').show()
print(spark.catalog.listTables())
print(spark.conf.get('spark.sql.warehouse.dir'))
```

```
+-----+
|databaseName|
+-----+
| default|
| project|
+-----+

++
||
++
++

+-----+-----+-----+
|database|tableName|isTemporary|
+-----+-----+-----+
| project| cars| false|
+-----+-----+-----+
```

Activate

```
+-----+-----+-----+-----+-----+-----+
|price|year|mileage|city|state|vin|make|model|
+-----+-----+-----+-----+-----+-----+
| null| null| null| City| State| Vin| Make| Model|
| 8995| 2014| 35725| EL Paso| TX| 19VDE2E53EE000083| Acura| ILX6-Speed|
| 10888| 2013| 19606| Long Island City| NY| 19VDE1F52DE012636| Acura| ILX5-Speed|
| 8995| 2013| 48851| EL Paso| TX| 19VDE2E52DE000025| Acura| ILX6-Speed|
| 10999| 2014| 39922| Windsor| CO| 19VDE1F71EE003817| Acura| ILX5-Speed|
| 14799| 2016| 22142| Lindon| UT| 19UDE2F32GA001284| Acura| ILXAutomatic|
| 7989| 2012| 105246| Miami| FL| JH4CU2F83CC019895| Acura| TSXAutomatic|
| 14490| 2014| 34832| Greatneck| NY| JH4CU2F84EC002686| Acura| TSXSpecial|
| 13995| 2013| 32384| West Jordan| UT| JH4CU2F64DC006203| Acura| TSX5-Speed|
| 10495| 2013| 57596| Waterbury| CT| 19VDE2E58DE000234| Acura| ILX6-Speed|
| 9995| 2013| 63887| EL Paso| TX| 19VDE1F50DE010450| Acura| ILX5-Speed|
| 12921| 2012| 58550| Boise| ID| JH4CU2F44CC003220| Acura| TSXAutomatic|
| 12000| 2013| 40527| Long Island City| NY| 19VDE1F38DE020867| Acura| ILX5-Speed|
| 7750| 2009| 91980| San Antonio| TX| JH4CU26639C015787| Acura| TSX4dr|
| 17628| 2015| 13797| Fargo| ND| 19VDE1F38FE001240| Acura| ILX5-Speed|
| 13999| 2013| 35035| Santa Ana| CA| JH4CU2F4XDC000369| Acura| TSX5-Speed|
| 14995| 2014| 23454| Hackettstown| NJ| 19VDE1F31EE009243| Acura| ILX5-Speed|
| 14990| 2015| 23603| Freeport| NY| 19VDE1F3XFE007606| Acura| ILX5-Speed|
| 14590| 2010| 19250| Clearwater| FL| JH4CU2F6XAC041680| Acura| TSX4dr|
| 9500| 2011| 68289| Arcadia| FL| JH4CU2F62BC007928| Acura| TSX4dr|
+-----+-----+-----+-----+-----+-----+

only showing top 20 rows

[Table(name='cars', database='project', description=None, tableType='EXTERNAL', isTemporary=False)]
/user/hive/warehouse
```

```
In [4]: df=spark.read.table("cars")
```

```
In [5]: df.columns
```

```
Out[5]: ['price', 'year', 'mileage', 'city', 'state', 'vin', 'make', 'model']
```

```
In [6]: df.show()
```

price	year	mileage	city	state	vin	make	model
10888	2013	19606	Long Island City	NY	19VDE1F52DE012636	Acura	ILX5-Speed
8995	2013	48851	El Paso	TX	19VDE2E52DE000025	Acura	ILX6-Speed
10999	2014	39922	Windsor	CO	19VDE1F71EE003817	Acura	ILX5-Speed
14799	2016	22142	Lindon	UT	19UDE2F32GA001284	Acura	ILXAutomatic
7980	2012	105246	Miami	FL	JH4CU2F83CC019895	Acura	TSXAutomatic
14490	2014	34832	Greatneck	NY	JH4CU2F84EC002686	Acura	TSXSpecial
13995	2013	32384	West Jordan	UT	JH4CU2F64DC006203	Acura	TSX5-Speed
10495	2013	57596	Waterbury	CT	19VDE2E50DE000234	Acura	ILX6-Speed
9995	2013	63887	El Paso	TX	19VDE1F50DE010458	Acura	ILX5-Speed
12921	2012	58550	Boise	ID	JH4CU2F44CC003228	Acura	TSXAutomatic
12000	2013	40527	Long Island City	NY	19VDE1F38DE020867	Acura	ILX5-Speed
7750	2009	91980	San Antonio	TX	JH4CU26639C015787	Acura	TSX4dr
17628	2015	13797	Fargo	ND	19VDE1F38FE001248	Acura	ILX5-Speed
13999	2013	35035	Santa Ana	CA	JH4CU2F4XDC000369	Acura	TSX5-Speed
14995	2014	23454	Hackettstown	NJ	19VDE1F31EE009243	Acura	ILX5-Speed
14990	2015	23603	Freeport	NY	19VDE1F3XFE007606	Acura	ILX5-Speed
14590	2010	19250	Clearwater	FL	JH4CU2F6XAC041680	Acura	TSX4dr
9500	2011	68289	Arcadia	FL	JH4CU2F62BC007928	Acura	TSX4dr

Activate
Go to Setting

Basic preprocessing

```
In [11]: df.dropDuplicates().show()
```

price	year	mileage	make	model
13998	2010	41652	Acura	TSX4dr
16750	2015	57483	Acura	ILX5-Speed
16517	2015	48918	Acura	ILX5-Speed
19968	2016	31019	Acura	ILXAutomatic
17384	2013	35257	Acura	TLBase
14539	2012	59046	Acura	TL2WD
7995	2008	78692	Acura	RDX4WD
4988	2000	109624	Acura	TLAutomatic
22500	2014	35341	Acura	TL2WD
15000	2011	80442	Acura	RDXFWD
20190	2015	31235	Acura	TLXFWD
23499	2014	48194	Acura	TL5H-AWD
21800	2015	32649	Acura	TLXFWD
27999	2016	10398	Acura	TLXFWD
27749	2016	15429	Acura	TLXFWD
25656	2015	29778	Acura	TLXFWD
16734	2012	77424	Acura	TL2WD
33990	2016	7205	Acura	RDXFWD
24420	2013	56175	Acura	MDXwith
4499	2003	184707	Acura	TLAutomatic

only showing top 20 rows

2.2 Data cleaning and Preprocessing

Another approach is using Python on google collab

Out of all independent columns, the column Vin acts as an Id for the car models so we will be dropping it since there is no significance to it. Afterwards, duplicated values (if present) were checked, pandas has an inbuilt function which returns duplicate rows present in the dataframe.

```
[ ] 1 #Vin is unnecessary column so we are dropping it
    2 df.drop('Vin',axis=1,inplace=True)
    3 df.shape
```

(852122, 7)

```
[ ] 1 #checking duplicate data that might be present in the dataset
    2
    3 dfdup=df[df.duplicated()]
    4 dfdup
```

	Price	Year	Mileage	City	State	Make	Model
314	22000	2017	10	Chicago	IL	Acura	ILXAutomatic
1259	23566	2017	16	Larchmont	NY	Acura	ILXPremium
6258	36000	2018	5	Littleton	CO	Acura	RDXAWD
6356	33900	2017	4250	Salt Lake City	UT	Acura	RDXAWD
7180	38275	2018	5	Littleton	CO	Acura	RDXAWD
...
819060	36998	2014	35370	Houston	TX	Toyota	4Runner4x4
819874	32018	2017	6	Mechanicsville	VA	Toyota	Tacoma2WD
824387	16995	2017	15	San Antonio	TX	Volkswagen	Passat1.8T
825336	16996	2017	7787	San Antonio	TX	Volkswagen	Passat1.8T
851162	39000	2016	7	San Francisco	CA	Volvo	XC60FWD

554 rows × 7 columns

So, it could be seen that the dataset has 554 duplicated rows, it's typically a good idea to remove duplicate data points so the model can better generalize to the full dataset. The duplicated rows can be dropped as:

```
[ ] 1 #checking shape of dataset
    2 df.shape
```

```
(852122, 7)
```

```
[ ] 1 #removing duplicated rows from the dataset
    2 df.drop_duplicates(inplace=True)
    3
    4 #again checking for shape of the now data
    5 df.shape
```

```
(851568, 7)
```

Afterwards, missing values were checked. It came out that the dataset had no missing values present.

```
[ ] 1 #checking for missing values
    2
    3 df.isnull().sum()
```

```
Price      0
Year        0
Mileage     0
City        0
State       0
Make        0
Model       0
dtype: int64
```

```
[ ] 1 #checking all datatypes
    2
    3 df.dtypes
```

```
Price      int64
Year        int64
Mileage     int64
City        object
State       object
Make        object
Model       object
dtype: object
```

2.3 Exploratory Data Analysis

Exploratory Data Analysis refers to the critical process of performing initial investigations on data so as to discover patterns and to check summary statistics and graphical representations.

Now we will be checking the unique number of values that are present in the categorical columns

```
[ ] 1 #checking number of unique values in the columns
    2
    3 df.nunique()
```

```
Price      47124
Year        22
Mileage    158836
City       2553
State       59
Make        58
Model      2736
dtype: int64
```

Unique values in the columns

```
[ ] 1 #checking out different classes or labels present in the variables
    2 cat=df.select_dtypes(include=['object']).columns
    3 for i in df.loc[:,cat]:
    4     print(df[i].value_counts(),'\n\n') ## value_count counts all unique values
    5     ##dropping all values who has count < 50
    6     counts = df[i].value_counts()
    7     df = df[~df[i].isin(counts[counts < 50].index)]
```

```
NH      5522
NE      4924
IA      4768
NM      4582
ID      3444
HI      2911
DE      2399
MT      1912
RI      1818
ME      1797
```

To check the correlation between all the columns present(categorical and numerical), we will be Encoding the categorical columns using Label encoder from sklearn

```
[ ] 1 df_lbe=df.copy()

[ ] 1 cat
    Index(['City', 'State', 'Make', 'Model'], dtype='object')

[ ] 1 #converting categorical variables into numerical ones using LabelEncoder so that we can check the coorelation between all variables
    2
    3 from sklearn.preprocessing import LabelEncoder
    4
    5 for i in cat:
    6     df_lbe[i] = LabelEncoder().fit_transform(df_lbe[i])
```

Checking the correlation of all variables with each other

```
[ ] 1 #checking out correlation matrix
    2
    3 df_lbe.corr()
```

	Price	Year	Mileage	City	State	Make	Model
Price	1.000000	0.437095	-0.444979	-0.015443	0.027127	-0.079957	0.076785
Year	0.437095	1.000000	-0.765428	0.010485	-0.023711	0.023872	-0.014434
Mileage	-0.444979	-0.765428	1.000000	-0.013848	0.024843	-0.032335	0.045073
City	-0.015443	0.010485	-0.013848	1.000000	-0.048887	0.007954	-0.003859
State	0.027127	-0.023711	0.024843	-0.048887	1.000000	-0.003653	0.006630
Make	-0.079957	0.023872	-0.032335	0.007954	-0.003653	1.000000	0.030264
Model	0.076785	-0.014434	0.045073	-0.003859	0.006630	0.030264	1.000000



Checking the correlation of all variables only with respect to price

```
[ ] 1 # Find most important features relative to target Price
    2
    3 print("Find most important features relative to target")
    4 corr = df_lbe.corr()
    5 corr.sort_values(["Price"], ascending = False, inplace = True)
    6 print(corr.Price)
```


```
Find most important features relative to target
Price      1.000000
Year       0.437095
Model      0.076785
State      0.027127
City      -0.015443
Make      -0.079957
Mileage   -0.444979
Name: Price, dtype: float64
```

Checking the correlation graphically using heatmap

```

1 # Price correlation matrix
2 k = 7 #number of variables for heatmap
3 corrmat = df_lbe.corr()
4 cols = corrmat.nlargest(k, 'Price')['Price'].index
5 cm = np.corrcoef(df_lbe[cols].values.T)
6 sns.set(font_scale=1.25)
7 hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_kws={'size': 10}, yticklabels=cols.values, xticklabels=cols.values)
8 plt.show()

```



Conclusion from correlation matrix:

Column 'Year' has a mild positive correlation with target 'Price'

Column 'Mileage' has a mild negative correlation with 'Price'

Columns 'Make' and 'Model' have low positive correlation with target 'Price'

Columns 'City' and 'State' have trivial correlation with the target column 'Price'

2.4 Feature selection

First, we had dropped 'City' and 'State' from the dataset, because it must not contribute to the prediction of Prices also both these categorical variables have large number of unique classes present, so it would be a bit complex in converting these categorical variables into numerical ones.

```
✓ [28] 1 #Dropping City and State as they have minimal correlation for prediction of price
0s      2 df.drop(['City', 'State'], axis=1, inplace=True)
```

```
✓ [29] 1 cat=df.select_dtypes(include=['object']).columns
0s
```

```
✓ [30] 1 df.head()
0s
```

	Price	Year	Mileage	Make	Model
1	10888	2013	19606	Acura	ILX5-Speed
3	10999	2014	39922	Acura	ILX5-Speed
4	14799	2016	22142	Acura	ILXAutomatic
5	7989	2012	105246	Acura	TSXAutomatic

Now we will use Standard Scaler from sklearn to apply scaling on the columns Year and Mileage. This will help some models to perform better and give relevant predictions.

```
✓ [30] 1 df.head()
0s
```

	Price	Year	Mileage	Make	Model
1	10888	2013	19606	Acura	ILX5-Speed
3	10999	2014	39922	Acura	ILX5-Speed
4	14799	2016	22142	Acura	ILXAutomatic
5	7989	2012	105246	Acura	TSXAutomatic
6	14490	2014	34032	Acura	TSXSpecial

```
✓ [31] 1 from sklearn.preprocessing import StandardScaler          #scaling the data
0s      2 ##Apply scaling on Independent variables only so price predicted in end by the model does not come scaled
      3
      4 sc= StandardScaler()
      5 sc.fit(df.loc[:,['Year','Mileage']])
      6 ##fitting on only x(Independent variables)
      7
      8 df.loc[:,['Year','Mileage']] = sc.transform(df.loc[:,['Year','Mileage']])
```


Outliers

Z score is also called standard score. This score helps to understand if a data value is greater or smaller than mean and how far away it is from the mean.

More specifically, Z score tells how many standard deviations away a data point is from the mean. If the z score of a data point is more than 3 or -3, it indicates that the data point is quite different from the other data points. Such a data point can be an outlier. So we are only selecting the rows that have value in any of the columns greater than -3 and less than 3.

Outlier removal

```
✓ 0s [2] 1 #function to select only rows whose value is > -3 and value < 3
2
3 def outlier_removal_zscore(df , cont_columns):
4     for col in cont_columns:
5         print("Before removing outliers from col=",col)
6         print("Shape =",df.shape)
7         df = df.loc[(df[col] >= -3)&(df[col] <= 3),:]
8         print("After removing outliers from col=",col)
9         print("Shape =",df.shape)
10    return df
11
12 df = outlier_removal_zscore(df.copy(),['Year','Mileage'])
```

```
↳ Before removing outliers from col= Year
Shape = (818681, 5)
After removing outliers from col= Year
Shape = (804931, 5)
Before removing outliers from col= Mileage
Shape = (804931, 5)
After removing outliers from col= Mileage
Shape = (797354, 5)
```

We'll be using One hot encoded data while training the model

But we'll only be sampling 50000 rows from all of the cleaned and scaled data.

```
✓ [33] 1 #converting categorical variables into numerical ones using One Hot Encoder
13s 2
3 df = pd.get_dummies(df)

✓ [34] 1 df.shape
0s
(797354, 1277)

✓ [35] 1 df=df.sample(50000,random_state=7) #sampling only limited amount of data since there are too many OHE columns
5s

✓ [36] 1 df.shape
0s
(50000, 1277)
```

After One hot encoding we have 1277 columns.

Saving One hot encoded data into an empty csv file

```
[ ] 1 # Saving the scaled and one hot encoded data with no outliers to an empty csv file
2 ##DO NOT RUN AGAIN
3 #df.to_csv("/content/drive/MyDrive/pro - Shreya & Rinky/cars_ohe_sc.csv",index= False, header=True)
```

2.5 Splitting X and Y

```
✓ [40] 1 #seperating the independant and dependant variables
0s 2 x = df.drop('Price',axis=1)
3 y = df['Price']
4
5 x.shape,y.shape

((50000, 1276), (50000,))
```

After splitting the data in x and y, and applying train test split, we are applying PCA to select only the significant values that will affect our target variable

2.6 Train test split and PCA

```
✓ [41] 1 #using train test split  
0s    2 x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=777,test_size=.30)
```

PCA

```
✓ [42] 1 from sklearn.decomposition import PCA  
28s    2  
    3 #applying pca to reduce dimensions and setting n_components to 0.98  
    4 pca = PCA(n_components = 0.98)  
    5  
    6 x_train = pca.fit_transform(x_train)
```

```
✓ [43] 1 x_test = pca.transform(x_test)  
0s
```

```
✓ [44] 1 explained_variance = pca.explained_variance_ratio_ #eigen values/ total eigen values  
0s
```

```
✓ [45] 1 x_train.shape  
0s  
(35000, 676)
```

After applying PCA we have 35000 rows and 676 columns which are the only relevant values for our Price prediction

3. Building Machine Learning Models

Now we will train several Machine Learning models and compare their results. Later on, we will use evaluation metrics.

3.1 Lasso

▼ Lasso Regression

```
1 #Lasso
2
3 ls=Lasso()
4 ls.fit(x_train,y_train)
5 predtrain=ls.predict(x_train)
6 predls=ls.predict(x_test)
7
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_coordinate_descent.py:145:
model = cd_fast.enet_coordinate_descent(
```

```
[ ] 1 #checking error for Lasso
2
3 msels=mean_squared_error(y_test,predls)
4 maels=mean_absolute_error(y_test,predls)
5 r2ls=r2_score(y_test,predls)
6 print(msels)
7 print(maels)
8 print(r2ls)
```

```
28910370.352732252
3200.514511821493
0.7990412586745311
```

3.2 Ridge

▼ Ridge Regression

```
[ ] 1 #Ridge
    2
    3 lrr=Ridge()
    4 lrr.fit(x_train,y_train)
    5 predlrr=lrr.predict(x_test)
    6 print(r2_score(y_test,predlrr)*100)
    7
```

81.70848343584713

```
[ ] 1 #checking error for Ridge
    2
    3 mselrr=mean_squared_error(y_test,predlrr)
    4 maelrr=mean_absolute_error(y_test,predlrr)
    5 r2lrr=r2_score(y_test,predlrr)
    6 print(mselrr)
    7 print(maelrr)
    8 print(r2lrr)
```

26314581.52528624
2947.8487873351924
0.8170848343584713

3.3 Random Forest

▼ Random Forest Regressor

✓
5m



```
1 #RandomForest
2
3 rf=RandomForestRegressor()
4 rf.fit(x_train,y_train)
5 predrf=rf.predict(x_test)
6 print(r2_score(y_test,predrf)*100)
```

80.41570919614958

✓
0s

```
[55] 1 #checking error for RandomForest
2
3 mserf=mean_squared_error(y_test,predrf)
4 maerf=mean_absolute_error(y_test,predrf)
5 r2rf=r2_score(y_test,predrf)
6 print(mserf)
7 print(maerf)
8 print(r2rf)
```

28174395.22662692
2880.621816409524
0.8041570919614958

3.4 Catboost

▼ Catboost Regression

✓
0s [56] 1 X = df.drop('Price',axis=1)
2 Y = df['Price']

✓
0s [57] 1 # Catboost requires validation dataset as well so we are implementing train_valid_test_split method from a fast_ml library
2
3 from fast_ml.model_development import train_valid_test_split
4 X_train,Y_train,X_valid,Y_valid,X_test,Y_test=train_valid_test_split(df,target='Price',train_size=0.6,valid_size=0.2,test_size=0.2,random_state=777)
5
6 X_train.shape,Y_train.shape, X_valid.shape,Y_valid.shape,X_test.shape, Y_test.shape

((30000, 1276), (30000,), (10000, 1276), (10000,), (10000, 1276), (10000,))

✓
0s [58] 1 #catboost needs to be given pool of data hence Pool method
2 train_dataset = cb.Pool(X_train, Y_train)
3 eval_dataset = cb.Pool(X_valid, Y_valid)

✓
0s [59] 1 #Creating model
2 cbt = cb.CatBoostRegressor(loss_function='RMSE')

```

✓ [60] 1 #Running possibilities for the best model
3m      2 grid = {'iterations': [ 250, 300],
3          'learning_rate': [ 0.2,0.5,0.7],
4          'depth': [8, 10],
5          'l2_leaf_reg': [0.2, 0.4]}
        6 cbt.grid_search(grid, train_dataset)

```

Streaming output truncated to the last 5000 lines.

41:	learn: 6955.9700441	test: 7255.1079661	best: 7255.1079661 (41)	total: 834ms	remaining: 5.13s
42:	learn: 6921.9623581	test: 7220.5414385	best: 7220.5414385 (42)	total: 854ms	remaining: 5.1s
43:	learn: 6881.1427863	test: 7184.1292651	best: 7184.1292651 (43)	total: 874ms	remaining: 5.08s
44:	learn: 6846.7009618	test: 7137.4235909	best: 7137.4235909 (44)	total: 892ms	remaining: 5.06s
45:	learn: 6816.2292652	test: 7119.5837000	best: 7119.5837000 (45)	total: 922ms	remaining: 5.09s
46:	learn: 6787.5088556	test: 7099.0851610	best: 7099.0851610 (46)	total: 940ms	remaining: 5.06s
47:	learn: 6760.7678046	test: 7080.0754851	best: 7080.0754851 (47)	total: 964ms	remaining: 5.06s
48:	learn: 6729.8440281	test: 7054.9357216	best: 7054.9357216 (48)	total: 987ms	remaining: 5.05s
49:	learn: 6704.0447037	test: 7030.7931330	best: 7030.7931330 (49)	total: 1s	remaining: 5.02s
50:	learn: 6673.2559944	test: 7005.1140692	best: 7005.1140692 (50)	total: 1.02s	remaining: 4.97s

```

✓ [61] 1 #fit on best model
4s      2 cbt.fit(train_dataset,eval_set=eval_dataset,early_stopping_rounds=50,plot=True,use_best_model=True, silent=False)

```

201:	learn: 3723.0781970	test: 4320.5763834	best: 4320.5763834 (201)	total: 2.7s	remaining: 641ms
202:	learn: 3719.4464177	test: 4319.0156706	best: 4319.0156706 (202)	total: 2.71s	remaining: 628ms
203:	learn: 3716.8812838	test: 4318.3066445	best: 4318.3066445 (203)	total: 2.72s	remaining: 614ms
204:	learn: 3712.6483496	test: 4317.8308741	best: 4317.8308741 (204)	total: 2.74s	remaining: 601ms
205:	learn: 3708.2120752	test: 4303.7900889	best: 4303.7900889 (205)	total: 2.75s	remaining: 588ms
206:	learn: 3705.5522389	test: 4301.3395999	best: 4301.3395999 (206)	total: 2.76s	remaining: 574ms
207:	learn: 3698.3753700	test: 4299.1454469	best: 4299.1454469 (207)	total: 2.78s	remaining: 561ms
208:	learn: 3693.6419538	test: 4292.2102901	best: 4292.2102901 (208)	total: 2.79s	remaining: 548ms
209:	learn: 3691.0667249	test: 4291.1599904	best: 4291.1599904 (209)	total: 2.8s	remaining: 534ms
210:	learn: 3688.5140968	test: 4288.4916059	best: 4288.4916059 (210)	total: 2.82s	remaining: 521ms
211:	learn: 3686.2703221	test: 4285.9722831	best: 4285.9722831 (211)	total: 2.83s	remaining: 507ms
212:	learn: 3679.7828635	test: 4284.3174402	best: 4284.3174402 (212)	total: 2.84s	remaining: 494ms
213:	learn: 3677.4227806	test: 4282.3576595	best: 4282.3576595 (213)	total: 2.86s	remaining: 481ms
214:	learn: 3672.6907165	test: 4278.7113929	best: 4278.7113929 (214)	total: 2.87s	remaining: 467ms
215:	learn: 3667.2899862	test: 4277.8750415	best: 4277.8750415 (215)	total: 2.88s	remaining: 454ms
216:	learn: 3661.0243113	test: 4274.2241038	best: 4274.2241038 (216)	total: 2.9s	remaining: 441ms
217:	learn: 3658.7163561	test: 4272.0864194	best: 4272.0864194 (217)	total: 2.91s	remaining: 428ms

```

✓ [62] 1 #evaluation metrics
0s      2 predcbt = cbt.predict(X_test)
        3 msecbt = (mean_squared_error(Y_test, predcbt))
        4 maecbt = mean_absolute_error(Y_test, predcbt)
        5 r2cbt = r2_score(Y_test, predcbt)
        6
        7 print(msecbt)
        8 print(maecbt)
        9 print(r2cbt)

```

```

20457770.70380485
2596.579351178401
0.851610217534246

```

3.5 XGBoost

▼ XGBoost

✓
2m

```
[63] 1 #XGBoost  
      2  
      3 xgb=XGBRegressor()  
      4 xgb.fit(x_train,y_train)  
      5 predxgb=xgb.predict(x_test)  
      6 print(r2_score(y_test,predxgb)*100)
```

77.84917476352999

✓
0s

```
[64] 1 #checking error for GradientBoosting  
      2  
      3 msexgb=mean_squared_error(y_test,predxgb)  
      4 maexgb=mean_absolute_error(y_test,predxgb)  
      5 r2xgb=r2_score(y_test,predxgb)  
      6 print(msexgb)  
      7 print(maexgb)  
      8 print(r2xgb)
```

31866668.59979161
3692.6774275065104
0.7784917476352998

3.6 Neural Network

▼ NN

```

✓ [65] 1 import math
2a    2 import pandas as pd
      3 import tensorflow as tf
      4 import matplotlib.pyplot as plt
      5 from tensorflow.keras import Model
      6 from tensorflow.keras import Sequential
      7 from tensorflow.keras.optimizers import Adam
      8 from sklearn.preprocessing import StandardScaler
      9 from tensorflow.keras.layers import Dense, Dropout
     10 from sklearn.model_selection import train_test_split
     11 from tensorflow.keras.losses import MeanSquaredLogarithmicError
     12

```

```

✓ [66] 1 # define model
3a    2 model1 = Sequential()
      3 model1.add(Dense(128, activation='relu'))
      4 model1.add(Dense(64, activation='relu'))
      5 model1.add(Dense(32, activation='relu'))
      6 # model1.add(Dense(8, activation='relu'))
      7 model1.add(Dense(1, activation='linear'))
      8

```

```

✓ [67] 1 # loss function
0a    2 # msle = MeanSquaredLogarithmicError()
      3
      4 model1.compile(
      5     | loss='mse',
      6     | optimizer=Adam(learning_rate=0.005),
      7     | metrics=['mae']
      8 )

```

```

✓ [68] 1 # train the model
1m    2 history = model1.fit(
      3     | x_train,
      4     | y_train,
      5     | epochs=25,
      6     | batch_size=64,
      7     | validation_split=0.2
      8 )

```

```

Epoch 1/25
438/438 [=====] - 8s 5ms/step - loss: 137140752.0000 - mae: 6966.0171 - val_loss: 30577726.0000 - val_mae: 3406.1130
Epoch 2/25
438/438 [=====] - 3s 6ms/step - loss: 30799670.0000 - mae: 2990.4158 - val_loss: 22988992.0000 - val_mae: 2995.0996
Epoch 3/25
438/438 [=====] - 2s 4ms/step - loss: 24534156.0000 - mae: 2720.6313 - val_loss: 19042686.0000 - val_mae: 2615.3176
Epoch 4/25
438/438 [=====] - 2s 4ms/step - loss: 21824194.0000 - mae: 2598.6443 - val_loss: 17809506.0000 - val_mae: 2532.5300

```

✓
0s [69] 1 model1.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	163456
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 32)	2080
dense_3 (Dense)	(None, 1)	33

=====
Total params: 173,825
Trainable params: 173,825
Non-trainable params: 0
=====

✓
1s [70] 1 prednn= model1.predict(x_test)
2 prednn
3

469/469 [=====] - 1s 2ms/step
array([[25024.29],
 [16655.719],
 [30460.64],
 ...,
 [7312.421],
 [20297.074],
 [29923.604]], dtype=float32)

✓
1s [71] 1 msenn, maenn = model1.evaluate(x_test,y_test)
2 msenn, maenn
3 r2nn=r2_score(y_test,prednn)

469/469 [=====] - 1s 2ms/step - loss: 20280962.0000 - mae: 2464.4556

✓
0s [72] 1 print(msenn, maenn, r2nn)

20280962.0 2464.45556640625 0.8590250582285328

3.7 Best Model?

```
[73] 1 ##ALL models used
2 ##                               Lasso, Ridge, Random Forest, Catboost, XGBoost, NN
3 results = pd.DataFrame({'Mean Squared Error':[msels,mselrr,mserf,msecbt,msexgb,msenn],
4                           'Mean Absolute Error':[maels,maelrr,maerf,maecbt,maexgb,maenn],
5                           'Accuracy Score':[r2ls*100,r2lrr*100,r2rf*100,r2cbt*100,r2xgb*100,r2nn*100]},
6                           index=['Lasso','Ridge','RandomForestRegressor','CatBoostRegressor','XGBRegressor','NeuralNetworkRegressor'])
7 result_df = results.sort_values(by='Accuracy Score', ascending=False)
8 result_df
```

	Mean Squared Error	Mean Absolute Error	Accuracy Score
NeuralNetworkRegressor	2.028096e+07	2464.455566	85.902506
CatBoostRegressor	2.045777e+07	2596.579351	85.161022
Ridge	2.631458e+07	2947.848787	81.708483
RandomForestRegressor	2.817440e+07	2880.621816	80.415709
Lasso	2.891037e+07	3200.514512	79.904126
XGBRegressor	3.186667e+07	3692.677428	77.849175

3.8 Random forest in Pyspark

```
In [12]: from pyspark.ml.linalg import Vector
         from pyspark.ml.feature import VectorAssembler
```

```
In [13]: df.columns
```

```
Out[13]: ['price', 'year', 'mileage', 'make', 'model']
```

```
In [16]: # Preprocessing: StringIndexer for categorical labels
         from pyspark.ml.feature import StringIndexer

         stringIndexer1 = StringIndexer(inputCol="model", outputCol="model_label")

         model=stringIndexer1.fit(df)
         indexed1=model.transform(df)

         stringIndexer2 = StringIndexer(inputCol="make", outputCol="make_label")
         model2=stringIndexer2.fit(indexed1)
         indexed2=model2.transform(indexed1)
```

```
In [17]: from pyspark.ml.linalg import Vector
         from pyspark.ml.feature import VectorAssembler

         assembler = VectorAssembler(inputCols=['year', 'mileage', 'make_label', 'model_label'], outputCol="features")
```

```
In [18]: indexed2.show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
|price|year|mileage| make|      model|model_label|make_label|
+-----+-----+-----+-----+-----+-----+-----+
| 8995|2013|  48851|Acura|  ILX6-Speed|      1611.0|      22.0|
|13995|2013|  32384|Acura|   TSX5-Speed|       726.0|      22.0|
|10495|2013|  57596|Acura|  ILX6-Speed|      1611.0|      22.0|
|12921|2012|  58550|Acura|TSXAutomatic|       646.0|      22.0|
|17628|2015|  13797|Acura|  ILX5-Speed|       300.0|      22.0|
|13999|2013|  35035|Acura|   TSX5-Speed|       726.0|      22.0|
|14995|2014|  23454|Acura|  ILX5-Speed|       300.0|      22.0|
|14990|2015|  23603|Acura|  ILX5-Speed|       300.0|      22.0|
|14590|2010|  19250|Acura|   TSX4dr|       431.0|      22.0|
| 9500|2011|  68289|Acura|   TSX4dr|       431.0|      22.0|
|16994|2015|  23946|Acura|  ILX5-Speed|       300.0|      22.0|
|15499|2014|  27171|Acura|   TSX5-Speed|       726.0|      22.0|
|13499|2014|  35037|Acura|  ILX5-Speed|       300.0|      22.0|
|14999|2014|  17669|Acura|  ILX5-Speed|       300.0|      22.0|
|14500|2010|  25926|Acura|   TSX4dr|       431.0|      22.0|
|16000|2015|  30881|Acura|  ILX5-Speed|       300.0|      22.0|
|17419|2015|  15390|Acura|  ILX5-Speed|       300.0|      22.0|
|14999|2015|  27333|Acura|  ILX5-Speed|       300.0|      22.0|
|14999|2015|  28326|Acura|  ILX5-Speed|       300.0|      22.0|
|17000|2015|  24671|Acura|  ILX5-Speed|       300.0|      22.0|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
In [19]: output=assembler.transform(indexed2)
```

```
In [20]: output.show(truncate=False)
```

```
+-----+-----+-----+-----+-----+-----+-----+
|price|year|mileage|make |model      |model_label|make_label|features|
+-----+-----+-----+-----+-----+-----+-----+
|8995 |2013|48851  |Acura|ILX6-Speed|1611.0     |22.0      |[2013.0,48851.0,22.0,1611.0]|
|13995|2013|32384  |Acura|TSX5-Speed|726.0      |22.0      |[2013.0,32384.0,22.0,726.0]|
|10495|2013|57596  |Acura|ILX6-Speed|1611.0     |22.0      |[2013.0,57596.0,22.0,1611.0]|
|12921|2012|58550  |Acura|TSXAutomatic|646.0      |22.0      |[2012.0,58550.0,22.0,646.0]|
|17628|2015|13797  |Acura|ILX5-Speed|300.0      |22.0      |[2015.0,13797.0,22.0,300.0]|
|13999|2013|35035  |Acura|TSX5-Speed|726.0      |22.0      |[2013.0,35035.0,22.0,726.0]|
|14995|2014|23454  |Acura|ILX5-Speed|300.0      |22.0      |[2014.0,23454.0,22.0,300.0]|
|14990|2015|23603  |Acura|ILX5-Speed|300.0      |22.0      |[2015.0,23603.0,22.0,300.0]|
|14500|2010|19250  |Acura|TSX4dr   |431.0      |22.0      |[2010.0,19250.0,22.0,431.0]|
|9500 |2011|68289  |Acura|TSX4dr   |431.0      |22.0      |[2011.0,68289.0,22.0,431.0]|
|16994|2015|23946  |Acura|ILX5-Speed|300.0      |22.0      |[2015.0,23946.0,22.0,300.0]|
|15499|2014|27171  |Acura|TSX5-Speed|726.0      |22.0      |[2014.0,27171.0,22.0,726.0]|
|13499|2014|35037  |Acura|ILX5-Speed|300.0      |22.0      |[2014.0,35037.0,22.0,300.0]|
|14999|2014|17669  |Acura|ILX5-Speed|300.0      |22.0      |[2014.0,17669.0,22.0,300.0]|
|14500|2010|25926  |Acura|TSX4dr   |431.0      |22.0      |[2010.0,25926.0,22.0,431.0]|
|16000|2015|30881  |Acura|ILX5-Speed|300.0      |22.0      |[2015.0,30881.0,22.0,300.0]|
|17419|2015|15390  |Acura|ILX5-Speed|300.0      |22.0      |[2015.0,15390.0,22.0,300.0]|
|14999|2015|27333  |Acura|ILX5-Speed|300.0      |22.0      |[2015.0,27333.0,22.0,300.0]|
|14999|2015|28326  |Acura|ILX5-Speed|300.0      |22.0      |[2015.0,28326.0,22.0,300.0]|
|17000|2015|24671  |Acura|ILX5-Speed|300.0      |22.0      |[2015.0,24671.0,22.0,300.0]|
+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

Activate
Go to Settings

```
In [21]: train_data, test_data = output.randomSplit([0.8, 0.2], seed=123)
```

```
In [22]: # Create a RandomForestRegressor
from pyspark.ml.regression import RandomForestRegressor
rf_regressor = RandomForestRegressor(featuresCol="features", labelCol="price", numTrees=150, maxBins=3000)
```

```
In [23]: rf_model = rf_regressor.fit(train_data)
```

```
In [24]: predictions = rf_model.transform(test_data)
```

```
In [25]: from pyspark.ml.evaluation import RegressionEvaluator
evaluator = RegressionEvaluator(labelCol="price", predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(predictions)
print("Root Mean Squared Error (RMSE):", rmse)
```

Root Mean Squared Error (RMSE): 6353.793822417411

```
In [26]: # R-squared (R2)
r2 = evaluator.evaluate(predictions, {evaluator.metricName: "r2"})
print("R-squared (R2):", r2)
```

R-squared (R2): 0.7792687452839657

4. Transformation on Pyspark

We have applied narrow and wide transformations on Pyspark and saved those respective dataframes into hdfs in csv file format

```
# transformations
```

```
In [29]: df=spark.read.table("cars")
```

```
In [41]: df=df.dropna()
```

```
In [42]: df1=df.orderBy(df.city.asc(),df.price.asc())
```

```
In [44]: df1.show() #cities and price sorted in ascending order
```

price	year	mileage	city	state	vin	make	model
10339	2011	102333	AKRON	OH	WDDGF8BB2BR150111	Mercedes-Benz	C-Class4dr
15889	2011	92812	AKRON	OH	4JGBB8GB0BA681651	Mercedes-Benz	M-Class4MATIC
16999	2008	39493	AKRON	OH	4JGBB22EX8A421443	Mercedes-Benz	M-Class4WD
17889	2011	72832	AKRON	OH	WDDHF8HB4BA501126	Mercedes-Benz	E-Class4dr
23876	2013	39961	AKRON	OH	WDCGG8JB1DG016238	Mercedes-Benz	GLK-ClassGLK350
25229	2014	66378	AKRON	OH	WDCGG8JB4EG326238	Mercedes-Benz	GLK-ClassGLK350
26449	2004	25403	AKRON	OH	WDBSK75F24F071548	Mercedes-Benz	SL-ClassSL500
33959	2014	35461	AKRON	OH	4JGDA5HB1EA353653	Mercedes-Benz	M-ClassML350
35889	2015	29610	AKRON	OH	WDCGG8JB8FG378795	Mercedes-Benz	GLK-ClassGLK350
4995	2004	187317	ALEXANDRIA	VA	2HKYF18534H553810	Honda	PilotEX
5795	2005	180159	ALEXANDRIA	VA	1HGCM66565A061282	Honda	Accord
5995	2005	106412	ALEXANDRIA	VA	JM1BK343151250325	Mazda	Mazda35dr
6795	2006	54181	ALEXANDRIA	VA	YV1MS390862216864	Volvo	S402.4L
8995	2012	53308	ALEXANDRIA	VA	1FAHP3N28CL408507	Ford	Focus5dr
8995	2010	26345	ALEXANDRIA	VA	KNAFW6A30A5256595	Kia	Forte
8995	2012	83654	ALEXANDRIA	VA	1YVHZ8DH7C5M35879	Mazda	Mazda64dr
9995	2013	20871	ALEXANDRIA	VA	1FADP3K21DL231433	Ford	FocusHatchback
11795	2014	104369	ALEXANDRIA	VA	JM3KE2CY6E0350691	Mazda	CX-5Touring
11995	2015	27124	ALEXANDRIA	VA	3FADP4CJXFM216202	Ford	FiestaSedan
12990	2008	75210	ALEXANDRIA	LA	WBAWB33548P135043	BMW	3

only showing top 20 rows

```
In [56]: df1.coalesce(1).write.format("csv").mode('overwrite').save("/user/talentum/tfml")
```

```
In [43]: df2=df.groupBy("year").count()
```

```
In [47]: df2.sort(df.year.desc()).show() #yearly count of cars sold in descending order
```

```
+---+-----+
|year| count|
+---+-----+
|2018|  922|
|2017| 91608|
|2016|132136|
|2015|157516|
|2014|162432|
|2013| 74701|
|2012| 49764|
|2011| 39768|
|2010| 27539|
|2009| 19061|
|2008| 24713|
|2007| 21171|
|2006| 15079|
|2005| 11005|
|2004|  8117|
|2003|  5649|
|2002|  3800|
|2001|  2584|
|2000|  1933|
|1999|  1254|
+---+-----+
```

only showing top 20 rows

```
In [57]: df2.coalesce(1).write.format("csv").mode('overwrite').save("/user/talentum/tfm2")
```

```
In [50]: df3=df.groupBy("make").count()
df3.show()
```

```
+-----+-----+
|      make| count|
+-----+-----+
| Volkswagen| 23249|
| Oldsmobile|  122|
|      AM|  19|
|      Lexus| 20641|
|      Jaguar| 2200|
|      Saturn|  963|
|      FIAT| 1782|
| Maserati| 1047|
| Rolls-Royce|  92|
|      Scion| 3043|
|      Jeep| 40373|
| Mitsubishi| 4080|
|      Kia| 28636|
| Chevrolet|102268|
|      Volvo| 5106|
|      Hyundai| 35837|
|      Saab|  260|
|      Honda| 50193|
| INFINITI| 12258|
|      MINI| 4375|
+-----+-----+
```

only showing top 20 rows

```
In [58]: df3.coalesce(1).write.format("csv").mode('overwrite').save("/user/talentum/tfm3")
```

```
In [62]: df4=df.filter(df.city=='El Paso').groupBy("make").count()
```

```
In [63]: df4.show()
```

```
+-----+-----+
|      make|count|
+-----+-----+
|Volkswagen| 100|
|   Lexus  |  34|
|   Jaguar |   5|
|   Saturn |   5|
|    FIAT  |  12|
| Maserati |   1|
|   Scion  |  12|
|   Jeep   | 193|
|Mitsubishi|  15|
|    Kia   |  77|
|Chevrolet | 475|
|   Volvo  |   2|
| Hyundai | 138|
|   Honda  | 184|
| INFINITI |  31|
|   MINI   |  19|
|    Audi  |  18|
|    Ram   |  66|
| Cadillac |  56|
| Pontiac  |   6|
+-----+-----+
only showing top 20 rows
```

```
In [64]: df4.coalesce(1).write.format("csv").mode('overwrite').save("/user/talentum/tfm4")
```

```
talentum@talentum-virtual-machine:~$ hdfs dfs -ls
```

```
Found 5 items
```

```
drwxr-xr-x - talentum supergroup 0 2023-08-18 21:29 cars_all
drwxr-xr-x - talentum supergroup 0 2023-08-30 08:17 tfm1
drwxr-xr-x - talentum supergroup 0 2023-08-30 08:17 tfm2
drwxr-xr-x - talentum supergroup 0 2023-08-30 08:17 tfm3
drwxr-xr-x - talentum supergroup 0 2023-08-30 08:20 tfm4
```

```
talentum@talentum-virtual-machine:~$ hdfs dfs -ls -R
```

```
drwxr-xr-x - talentum supergroup 0 2023-08-18 21:29 cars_all
-rw-r--r-- 1 talentum supergroup 56287873 2023-08-18 21:29 cars_all/true_car_listings.csv
drwxr-xr-x - talentum supergroup 0 2023-08-30 08:17 tfm1
-rw-r--r-- 1 talentum supergroup 0 2023-08-30 08:17 tfm1/_SUCCESS
-rw-r--r-- 1 talentum supergroup 54597868 2023-08-30 08:17 tfm1/part-00000-8140de3a-1141-4d8b-bb6c-2ccc7d2624eb-c000.csv
drwxr-xr-x - talentum supergroup 0 2023-08-30 08:17 tfm2
-rw-r--r-- 1 talentum supergroup 0 2023-08-30 08:17 tfm2/_SUCCESS
-rw-r--r-- 1 talentum supergroup 696 2023-08-30 08:17 tfm2/part-00000-5d8f28b0-91cc-4acc-aa21-9dc6f576dc24-c000.csv
drwxr-xr-x - talentum supergroup 0 2023-08-30 08:17 tfm3
-rw-r--r-- 1 talentum supergroup 0 2023-08-30 08:17 tfm3/_SUCCESS
-rw-r--r-- 1 talentum supergroup 696 2023-08-30 08:17 tfm3/part-00000-9be698f4-de65-4988-bea9-3930909e8d7c-c000.csv
drwxr-xr-x - talentum supergroup 0 2023-08-30 08:20 tfm4
-rw-r--r-- 1 talentum supergroup 0 2023-08-30 08:20 tfm4/_SUCCESS
-rw-r--r-- 1 talentum supergroup 380 2023-08-30 08:20 tfm4/part-00000-d7f80cd5-fb65-4e6c-b291-113d0bf9fcb2-c000.csv
```

```
talentum@talentum-virtual-machine:~$ hdfs dfs -ls tfm3
Found 2 items
-rw-r--r--  1 talentum supergroup          0 2023-08-30 08:17 tfm3/_SUCCESS
-rw-r--r--  1 talentum supergroup    696 2023-08-30 08:17 tfm3/part-00000-9be698f4-de65-4988-bea9-3930909e8d7c-c000.csv
talentum@talentum-virtual-machine:~$ hdfs dfs -cat tfm3/part-00000-9be698f4-de65-4988-bea9-3930909e8d7c-c000.csv
Volkswagen,23249
Oldsmobile,122
AM,19
Lexus,20641
Jaguar,2200
Saturn,963
FIAT,1782
Maserati,1047
Rolls-Royce,92
Scion,3043
Jeep,40373
Mitsubishi,4080
Kia,28636
Chevrolet,102268
Volvo,5106
Hyundai,35837
Saab,260
Honda,50193
INFINITI,12258
MINI,4375
Audi,12618
Lamborghini,121
Ram,19808
Cadillac,15047
Isuzu,76
Plymouth,51
```

Activate
Go to Settings

5. Analysis on Tableau

The analysis of data was performed on

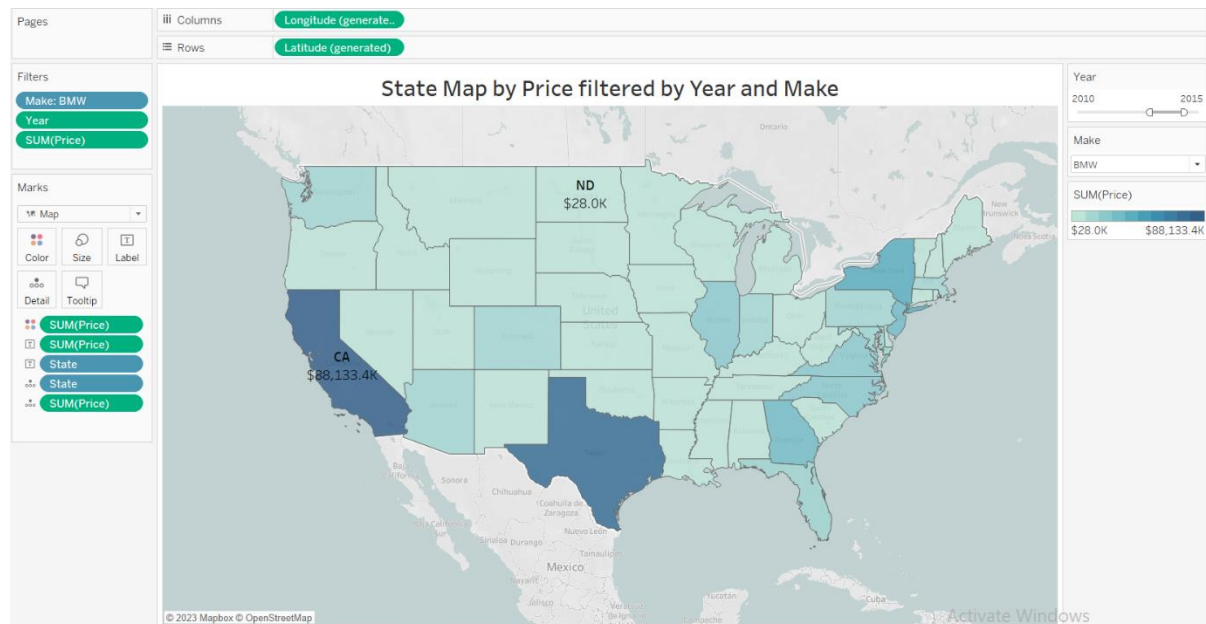


Fig 5.1 State map by price filtered by Year and Make




Fig 5.2 Top 10 Models by Price bar chart




Fig 5.3 Bottom 10 Models by Price bar chart




Fig 5.4 Top 10 Make by Price bar chart




Fig 5.5 Bottom 10 Make by Price bar chart




Fig 5.6 Dashboard of Top and Bottom features

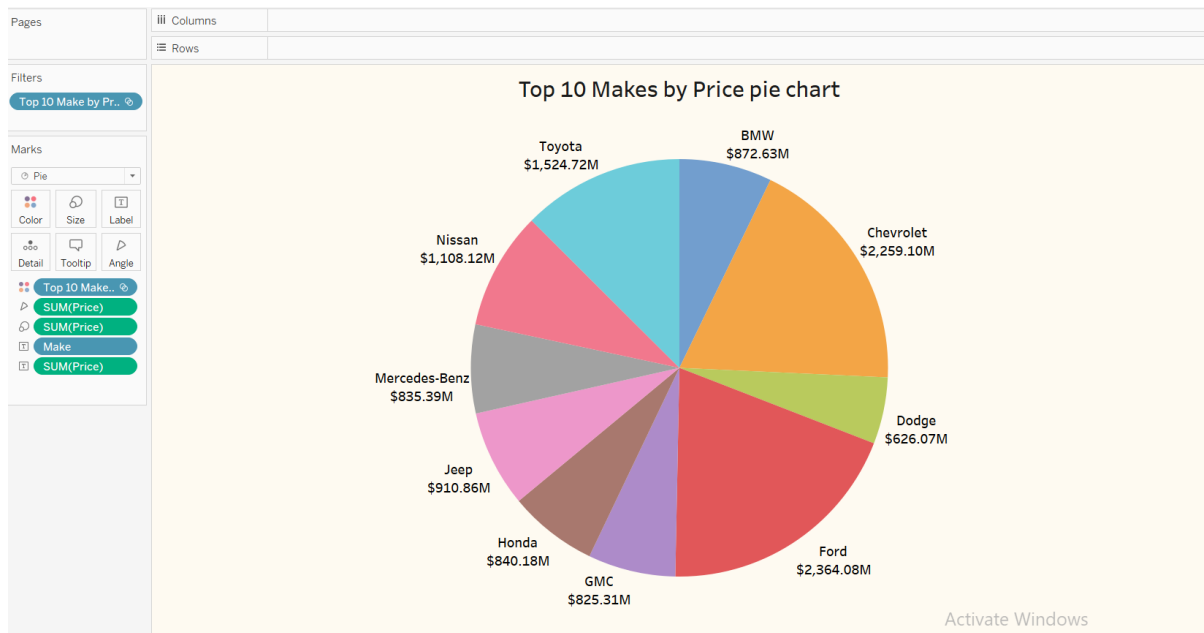


Fig 5.7 Top 10 Makes by Price pie chart




Fig 5.8 Price by Mileage scatter plot




Fig 5.9 Price by Year bar chart

6. Conclusion

The objective of this project is to fit models to predict the used cars price and find some important aspects of the used cars. In order to achieve this goal, we have cleaned and preprocessed the data, applied EDA and Feature selection, PCA for dimensionality reduction and then fit five different regression models to the dataset: lasso regression, ridge regression, random forest, catboost, xgboost and neural network.

The best two models were:

Neural network performed the best amongst all models with 3 Dense layers and activation function as RELu. Catboost had dataset split into train, validation and split parts . With further fine tuning there is a chance to find the best possible parameters for regression and improve the models further.

We have also performed ETL using big data technologies and performed transformation on Pyspark and stored them to HDFS.

In addition to Prediction of Prices we have also performed Analysis on our data on Tableau.

7. References

- <https://www.geeksforgeeks.org/principal-component-analysis-with-python/>
- <https://towardsdatascience.com/catboost-regression-in-6-minutes-3487f3e5b329>
- <https://www.analyticsvidhya.com/blog/2021/08/a-walk-through-of-regression-analysis-using-artificial-neural-networks-in-tensorflow/>
- <https://medium.com/vedity/python-data-preprocessing-using-pyspark-cc3f709c3c23>
- <https://www.machinelearningplus.com/pyspark/pyspark-random-forest/>